

# Toward Intelligent Political Actors

David Masad

April 28, 2014

## 1 Introduction

The model presented by Axelrod (1997), commonly known as the Tribute Model, is an early example of what Epstein and Axtell (1996) refer to as ‘artificial societies’ – agent-based models in which simple rules lead to the emergence of phenomena with recognizable real-world analogue. Agents in the Tribute Model represent atomic political actors, interacting with each other via demands for tribute (hence the model’s name) which are responded to with either payment or war. Past interactions may yield future cooperation, leading to the emergence of stable blocs of agents acting together against rival blocs. A similar approach was used by Axelrod’s student Lars-Erik Cederman in the GeoSim family of models (Cederman, 2003), also with rich results.

Both Axelrod and Cederman’s agents use simple heuristics to decide whether to go to war. While it is indeed an interesting result that many features of the international system emerge from agents using such heuristics, this is also a major departure of the models from both empirical evidence and qualitative theory. Even ancient proto-states carefully collected information and evaluated potential courses of action before going to war (Sheldon, 1989), and contemporary states have entire departments devoted to foreign policy planning. The Realist school of international relations theory goes even further, arguing that states make decisions rationally in response to the entirety of the information available to them (Waltz, 2010).

In this paper, I introduce a simplified model based primarily on the Axelrod Tribute model. I remove the commitment feature of the original model, and replace the linear topology with a more complex network. Thus, the model consists of atomic agents competing with their neighbors on a network.

I run the model using the original heuristics, showing that it display similar dynamics to the original. I then endow the agents with an ability to look ahead and anticipate the results of different actions and other agents’ responses to them. I explore whether endowing the agents with this expanded intelligence changes the dynamics of the model, and whether agents capable of looking ahead further than others are more successful.

## 2 Model Description

### 2.1 Model Architecture

The model consists of **agents**, endowed with **wealth**; agents are nodes in an undirected graph which governs their interactions. Each step of the model, one of the agents is activated. The active agent may then choose one of its graph neighbors to threaten (though it may also choose to threaten none).

A threatened agent now chooses whether to pay tribute or go to war. Paying tribute involves transferring a fixed, constant amount of wealth from the threatened agent to the threatening agent. War involves each agent reducing the opponent’s wealth by a fixed fraction of their own wealth. In any case, an agent’s wealth may not fall below zero.

After a set number of agents have been activated, the wealth of all agents increases by a fixed amount (referred to by Axelrod at the ‘harvest’).

In the original model, agent activation was random. In order to reduce the search space here, I make the agent activation order a known, fixed cycle, such that an agent will not be activated again before every other agent has been activated.

Parameter	Default value
Number of agents	10
Network topology	Small-World
War cost	0.25
Tribute cost	100
Starting wealth range	[300, 500]
Wealth growth per step	20

Table 1: Model parameters

### 2.2 Agent Decisionmaking

The model architecture as described above represents the external environment within which the agents function, and remains fixed over the course of this project. The purpose of the study, however, is to analyze the effect of agent decisionmaking – the internal environment. To do so, I implement three decisionmaking rules: the original heuristics described by Axelrod, relative wealth maximization with no lookahead, and relative wealth maximization with lookahead. In all cases, agents are endowed with perfect information. I will then analyze how agent behavior, and model outcomes, vary with each decision model.

#### 2.2.1 Heuristic Decisionmaking

As mentioned above, Axelrod’s original model involves very simple decision-making heuristics.

**Attacker** The active agent evaluates the *vulnerability* of each potential target (in this case, network neighbor). Agent  $i$  evaluates agent  $j$ 's vulnerability as follows:

$$vulnerability_i(j) = \frac{wealth_i - wealth_j}{wealth_j}$$

Agent  $i$  then chooses to threaten the neighbor with the greatest vulnerability score. If no neighbor has vulnerability greater than 0, no threat is issued.

**Defender** The threatened agent makes a simple cost-benefit analysis. If the tribute cost is less than the damage that would be incurred in a war with the attacker, the tribute is paid. Otherwise, the agent goes to war.

### 2.2.2 Relative Wealth Maximization

While the above heuristics are sufficient to yield the emergence of interesting phenomena in the original model, they bear only a cursory resemblance to qualitative theories of political decisionmaking. International relations theory argues that all power is relative, and must be measured against the resources of potential adversaries. Thus, for this behavior and the following one, I endow the agents with the objective of maximizing their wealth relative to that of the wealthiest neighbor.

**Defender** Under this decisionmaking rule, threatened agents evaluate their relative wealth immediately after paying tribute or going to war, and then choose the course of action which maximizes this value. Note that this allows an agent to choose war even if its cost exceeds the tribute payment, if the opponent is the wealthiest neighbor and war will improve their relative wealth.

**Attacker** For each possible target, the active agent determines that target's response to a threat, using the rule described above. The agent will then only issue a threat if the result of the target's decision will improve the active agent's own relative wealth.

### 2.2.3 Relative Wealth Maximization with Lookahead

Here is this project's novel addition to previous work: under this rule, agents forecast the state of the model several steps forward under each possible course of action (choice of target for the attacker, and decision between tribute or war for the defender). This allows the agents to take into account not only the immediate consequence of their decisions, but the second- or third-order consequences. For example, an agent may assess that by paying tribute, they will make the receiver a more tempting target for one of *its* neighbors, leading to a war between them in the next turn and thus improving the agent's own position.

In order to simulate the model forward, the agents are in fact simulating the decisionmaking process of the other agents. In order to avoid infinite recursion, there is a maximum depth limit imposed; when that depth limit is reached, the agents being simulated at that depth use the static Relative Wealth Maximization decision rule described above.

In practice, lookahead is implemented by spawning a copy of the model object, which includes copies of all the agent objects within it. A possible course of action is executed within the model copy, and then its step method is called, executing the next step. At each level, agents choose courses of action by recursively copying the model object to which they belong, until the depth limit is reached.

In order to both impose realism and limit the search space, agents are rendered incapable of simulating agents with a maximum recursion depth greater than their own. Thus, for example, an agent with recursion depth of 5 will simulate an agent with recursion depth 3 perfectly, but the depth-3 agent will simulate the depth-5 agent as though it were a depth-3 as well.

This depth limit has several interesting implications. Key among them is that unlike in the static rule above, the attacker’s simulation of the defender’s decisionmaking process is not identical to the defender’s actual decision. This introduces the opportunity for miscalculation without adding artificial noise to the decision process, corresponding to the political science literature arguing both that states are rational actors and that wars (particularly large wars) often arise due to strategic miscalculation of a rival’s response.

The addition of lookahead presents several opportunities for experiments. By varying the depth and horizon of agent forecasting, I test whether these have an effect on the frequency or magnitude of wars in the model. Furthermore, by endowing agents with heterogeneous recursion depth, I test the effect that increased ‘intelligence’ (both in the computational and political sense) has on outcome, as well as whether depth or horizon has a stronger effect.

### 3 Results

In order to investigate the model, I run it multiple times in four configurations, as follows:

1. Agents use heuristic decisionmaking
2. Agents use relative wealth maximization, with no lookahead
3. Agents use wealth maximization and have heterogenous lookahead:
  - (a) Variable agent graph
  - (b) Fixed agent graph

All configurations are implemented with the parameters described above, and were run for 300 steps, or 100 turns. For each, I recorded the wealth series for each agent, as well as the resources expended in wars. I also store the

network, as well as each agent’s initial wealth and (if relevant) their lookahead depth.

### 3.1 Heuristic Decisionmaking

When agents are given the heuristic decisionmaking rule, the results qualitatively align with the original Axelrod model. Figure X shows the wealth of all agents over the course of a sample run, demonstrating a clear divergence into high-wealth and low-wealth agents. This divergence occurs across all the model runs, as shown in Figure X.

[[ FIGURES HERE]]

[[ MORE ROBUST ANALYSIS OF WARS HERE]]

### 3.2 Relative Wealth Maximization

Figure X shows the distribution of agent wealth across all 50 of the experiment 2 models. Like the previous distribution, it is bimodal; however, it is otherwise clearly qualitatively different, with an approximately symmetric distribution around the higher-valued peak. While there may still be a core and periphery structure, the cores are more differentiated from one another.

[[ FIGURES HERE]]

The distribution of war frequencies is also different, with war-free turns significantly more likely than under the heuristic. This is perhaps the clearest effect of the rational decisionmaking rule – agents will only begin wars if they will lead to a certain, immediate relative gain, and the model confirms the intuition that they will thus engage in fewer wars overall.

### 3.3 Relative Wealth Maximization with Lookahead

The two final experiments are similar in that they involve agents with heterogeneous lookahead depths. While with experiment 3 a new interaction graph is generated for every iteration of the model, experiment 4 involves a fixed interaction graph, though the agents’ initial conditions and activation sequence are randomized with every iteration. Qualitatively, both result in wealth distributions that resemble experiment 2. Note, however, that experiment 4 appears to have a tighter distribution around the upper mode, as well as a larger number of wars.

#### 3.3.1 Heterogenous Depths

Greater recursion depth corresponds to increased intelligence, in that it allows agents to more accurately predict the decisions of the other agents further ahead. This is particularly true if agents in a model instance have heterogenous recursion depths, allowing high-depth agents to fully anticipate low-depth agents, who cannot anticipate the other agents in return. Thus, it seems reasonably to

expect that when agent recursion depths are heterogenous, high-depth agents will have greater success (and achieve higher wealth) than low-depth agents.

In order to test this hypothesis, I analyze the results of experiments 3 and 4 on an agent-by-agent basis. In addition to each agent’s lookahead depth, I examine the effects of other potential factors: the agents’ initial wealth, and their network degree.

I estimate these effects using a linear regression; while I have no reason to believe that the relationship is necessarily linear, this methodology should nevertheless be able to estimate the direction, magnitude and significance of the effects of each of the above characteristics on final agent wealth.

Tables X and X show the results of these regressions for experiments 1 and 2. Both indicate similar results: the degree of an agent has a strong negative effect on its final wealth, while starting wealth has a small positive effect on it. The recursion depth, however, has essentially no effect: in both regressions, the coefficient is small and statistically insignificant.

## 4 Discussion

The model presented here is an extremely simplified model of an international system. Unlike many previous models, its agents do not operate based on heuristics but robust adversarial search, choosing actions which maximize their relative wealth. This decision model does result in qualitative changes to the model dynamics, reducing the number and frequency of wars and widening the distribution of agent outcomes. Nevertheless, it appears that degree of each agent’s intelligence – the depth to which it can look ahead – does not affect its final wealth. This is an important result, for several reasons. It provides some validation for the use of simple decisionmaking heuristics, as the model behavior does not significantly diverge between the heuristic and lookahead decision rules.

The original Axelrod model, and most other models of international systems, involve random agent activation. Implementing lookahead taking random activation into account adds several additional layers to the decisionmaking process: it would require agents to either account for all possible activation order (massively increasing the search space), or else to stochastically sample possible activations. Furthermore, facing ranges of possible outcomes, agents must weight risks and rewards, introducing implicit or explicit risk-tolerance, which may also vary between agents. However, if increased lookahead when future activations are perfectly known does not increase agent outcomes, it is unlikely to do so in the stochastic case. Instead, the stochastic elements – both in actual activation and internal agent forecasting – would simply introduce noise.

The computational resources required to implement lookahead are significant, as every agent is simulating the responses of their neighbors, those simulated agents are simulating the actions of their own neighbors, and so forth until the recursion limit is reached. The time required thus increases exponentially with each additional level of depth added. Future work may need to develop applicable pruning methods, potentially including a ‘horizon’ on the network

beyond which agents will have a minimal effect on the currently-active agent, and thus not need to be simulated. This is important as the computational cost is what prevented this analysis from going beyond a depth of 5; it is possible that there is a qualitative shift in the effect of lookahead above a certain horizon, which has not been reached here.

## References

- Axelrod, R. M. (1997). *The Complexity of Cooperation: Agent-Based Models of Competition and Collaboration*. Princeton University Press.
- Cederman, L.-E. (2003). Modeling the size of wars: from billiard balls to sand-piles. *American Political Science Review* 97(01), 135-150.
- Epstein, J. M. and R. L. Axtell (1996, October). *Growing Artificial Societies: Social Science from the Bottom Up* (First ed.). A Bradford Book.
- Sheldon, R. M. (1989). Spying in mesopotamia: The world's oldest classified documents. *Studies in Intelligence* 33(1), 712.
- Waltz, K. N. (2010). *Theory of international politics*. Waveland Press.